

Operating Systems

Lecture 15

Reliable fs

Prof. Mengwei Xu

Threads to FS Reliability

- Operation interruption
 - A crash or power failure
 - A file operation often consists of many I/O updates to the storage
 - An example: `mv ./dir1/file1 ./dir2/file2`

Threads to FS Reliability

- Operation interruption
 - A crash or power failure
 - A file operation often consists of many I/O updates to the storage
 - An example: `mv ./dir1/file1 ./dir2/file2`
 - ❑ Writing the `dir1` directory file to remove `file1`
 - ❑ Growing the `dir2` directory's file to include another block of storage to accommodate a new directory entry for `file2`
 - ❑ Writing the new directory entry to the directory file
 - ❑ Updating the last-modified time of the `dir1` directory
 - ❑ Updating the file system's free space bitmap
 - ❑ Updating the size and last-modified time of the `dir2` directory
 - At a physical level, operations complete one at a time

Threads to FS Reliability

- Operation interruption
 - A crash or power failure
 - A file operation often consists of many I/O updates to the storage
 - An example: `mv ./dir1/file1 ./dir2/file2`
 - At a physical level, operations complete one at a time
- Loss of stored data
 - Either physical or electric

What a Reliable FS Does?

- “All or nothing”
 - Either an update is completed, or not at all
 - Must be guaranteed whenever a crash happens
 - Must be transparent to users/apps
 - An example: transfer \$100 from Bob’s account to Alice’s account
- Quite similar to the critical section problem in concurrency
 - Avoid someone observing the state in an intermediate, inconsistent state
 - No control over “when it happens”

Goals for Today

- Transactions for atomic updates
 - Redo Logging
- Redundancy for media failures
 - RAID

Goals for Today

- Transactions for atomic updates
 - Redo Logging
- Redundancy for media failures
 - RAID

Reliability Approach #1: Careful Ordering

- Sequence operations in a specific order
 - Careful design to allow sequence to be interrupted safely
- Post-crash recovery
 - Read data structures to see if there were any operations in progress
 - Clean up/finish as needed
- Approach taken by
 - FAT and FFS (**fsck**) to protect filesystem structure/metadata
 - Many app-level recovery schemes (e.g., Word, emacs autosaves)

FFS: Create a File

Normal operation:

- Allocate data block
- Write data block
- Allocate inode
- Write inode block
- Update bitmap of free blocks and inodes
- Update directory with file name → inode number
- Update modify time for directory

Recovery (file system check, `fsck`) :

- Scan inode table
- If any unlinked files (not in any directory), delete or put in lost & found dir
- Compare free block bitmap against inode trees
- Scan directories for missing update/access times

Time proportional to disk size

Issues with Approach #1

- Complex reasoning
 - So many possible operations and failures
- Slow updates
 - File systems are forced to insert sync operations or barriers between dependent operations
- Extremely slow recovery
 - Need to scan all of its disks for inconsistent metadata structures

Transactions

- Use *Transactions* (事务) for atomic updates
 - Ensure that multiple related updates are performed atomically
 - i.e., if a crash occurs in the middle, the state of the systems reflects either *all or none* of the updates
 - Most modern file systems use transactions internally to update filesystem structures and metadata
 - Many applications implement their own transactions
- They extend concept of atomic update from memory to stable storage
 - Atomically update multiple persistent data structures

Transactions

- An **atomic sequence** of actions (reads/writes) on a storage system (or database)
- That takes it from one **consistent state** to another



Typical Structure

- **Begin** a transaction – get transaction id
- Do a bunch of updates
 - If any fail along the way, **roll-back**
 - Or, if any conflicts with other transactions, **roll-back**
- **Commit** the transaction

“Classic” Example: Transaction

```
BEGIN;      --BEGIN TRANSACTION
```

```
UPDATE accounts SET balance = balance - 100.00 WHERE  
    name = 'Alice';
```

```
UPDATE branches SET balance = balance - 100.00 WHERE  
    name = (SELECT branch_name FROM accounts WHERE name =  
    'Alice');
```

```
UPDATE accounts SET balance = balance + 100.00 WHERE  
    name = 'Bob';
```

```
UPDATE branches SET balance = balance + 100.00 WHERE  
    name = (SELECT branch_name FROM accounts WHERE name =  
    'Bob');
```

```
COMMIT;      --COMMIT WORK
```

Transfer \$100 from Alice's account to Bob's account

The Key Properties of Transactions

- **Atomicity:** all actions in the transaction happen, or none happen
- **Consistency:** transactions maintain data integrity, e.g.,
 - Balance cannot be negative
 - Cannot reschedule meeting on February 30
- **Isolation:** execution of one transaction is isolated from that of all others; no problems from concurrency
- **Durability:** if a transaction commits, its effects persist despite crashes

Logging

- Instead of modifying data structures on disk directly, write changes to a journal/log
 - Intention list: set of changes we intend to make
 - Log/Journal is **append-only**
 - Single commit record commits transaction
- Once changes are in log, it is safe to apply changes to data structures on disk
 - Recovery can read log to see what changes were intended
 - Can take our time making the changes
 - ❑ As long as new requests consult the log first
- Basic assumption:
 - Updates to sectors are atomic and ordered

Implementing Transactions: Redo Logging

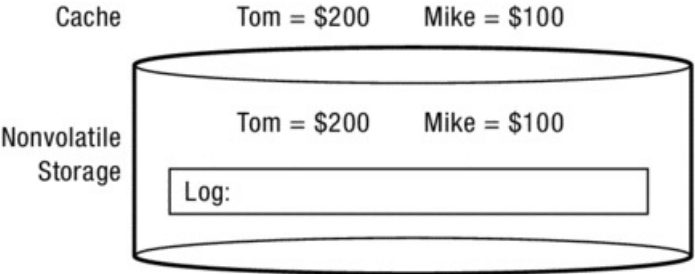
- Prepare
 - Write all changes/updates to log (日志)
 - Can happen at once, or over time
 - Wait until all updates are written in log
- Commit
 - Append a commit record to the log
 - Or can roll back, write a roll-back record
- Write-back
 - Write all of the transaction's updates to disk
- Garbage collection
 - Reclaim space in log
- Recovery
 - Read log
 - Redo any operations for committed transactions
 - Garbage collect log

Implementing Transactions: Redo Logging

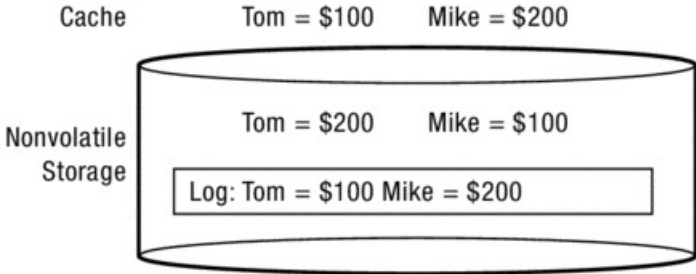
- Prepare
 - Write all changes/updates to log (日志)
 - Can happen at once, or over time
 - Wait until all updates are written in log
 - Commit
 - Append a commit record to the log
 - Or can roll back, write a roll-back record
 - Write-back
 - Write all of the transaction's updates to disk
 - Garbage collection
 - Reclaim space in log
- An atomic operation
- Before it, we can safely roll-back
 - After it, the transaction must take effect

Example #1

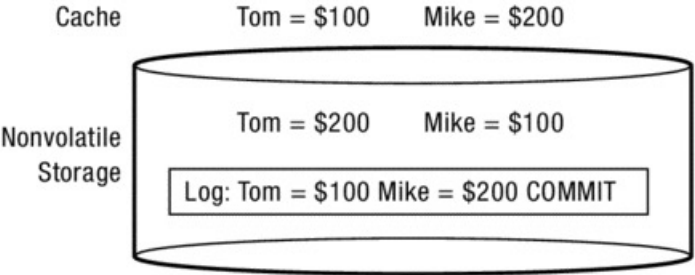
a) Original state



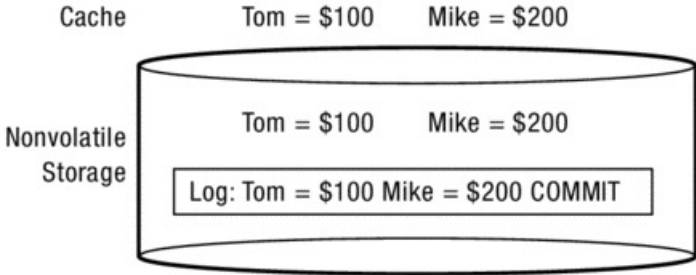
b) Updates appended to log



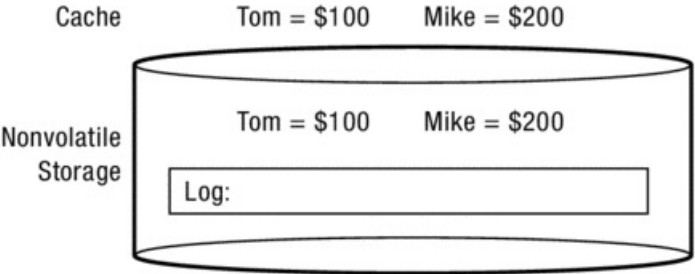
c) Commit appended to log



d) Updates applied




e) Garbage collect completed transactions from log



Example #2: Creating a File

- Find free data block(s)
- Find free inode entry
- Find dirent insertion point


-
- Write map (i.e., mark used)
 - Write inode entry to point to block(s)
 - Write dirent to point to inode



Example #2: Creating a File


- Find free data block(s)
- Find free inode entry
- Find dirent insertion point

-
- [log] Write map (used)
 - [log] Write inode entry to point to block(s)
 - [log] Write dirent to point to inode




ReDo Log

- After Commit
- All access to file system first looks in log
- Eventually copy changes to disk




Crash During Logging – Recover

- Upon recovery scan the log
- Detect transaction start with no commit
- Discard log entries
- Disk remains unchanged



Recovery After Commit

- Scan log, find start
- Find matching commit
- Redo it as usual
 - Or just let it happen later



Implementation Details

- Deal with concurrent transactions
 - Must identify which transaction does a record belong to
- Repeated write-backs are OK
 - Works for idempotent (幂等) updates: “write 42 to each byte of sector 74”
 - Redo log systems do not permit non-idempotent records such as “add 42 to each byte in sector 74”.
- Restarting recovery is OK
 - If another crash occurs during recovery

Implementation Details

- The performance of redo logging is not as bad as it looks like:
 - Log updates are sequential
 - Asynchronous write-back
 - ❑ Low latency for commit(); high throughput as updates can be batched
 - Group commit: combine a set of transaction commits into one log write
 - ❑ Amortize the cost of initiating the write (e.g., seek and rotational delays).
- New requests (e.g., reads) need to consult the log first to ensure the data consistency
 - Can be alleviated by caching
- Ordering is essential, as we must ensure:
 - A transaction's updates are on disk in the log before the commit is
 - The commit is on disk before any of the write-backs are
 - All of the write-backs are on disk before a transaction's log records are garbage collected.

Transactional File Systems

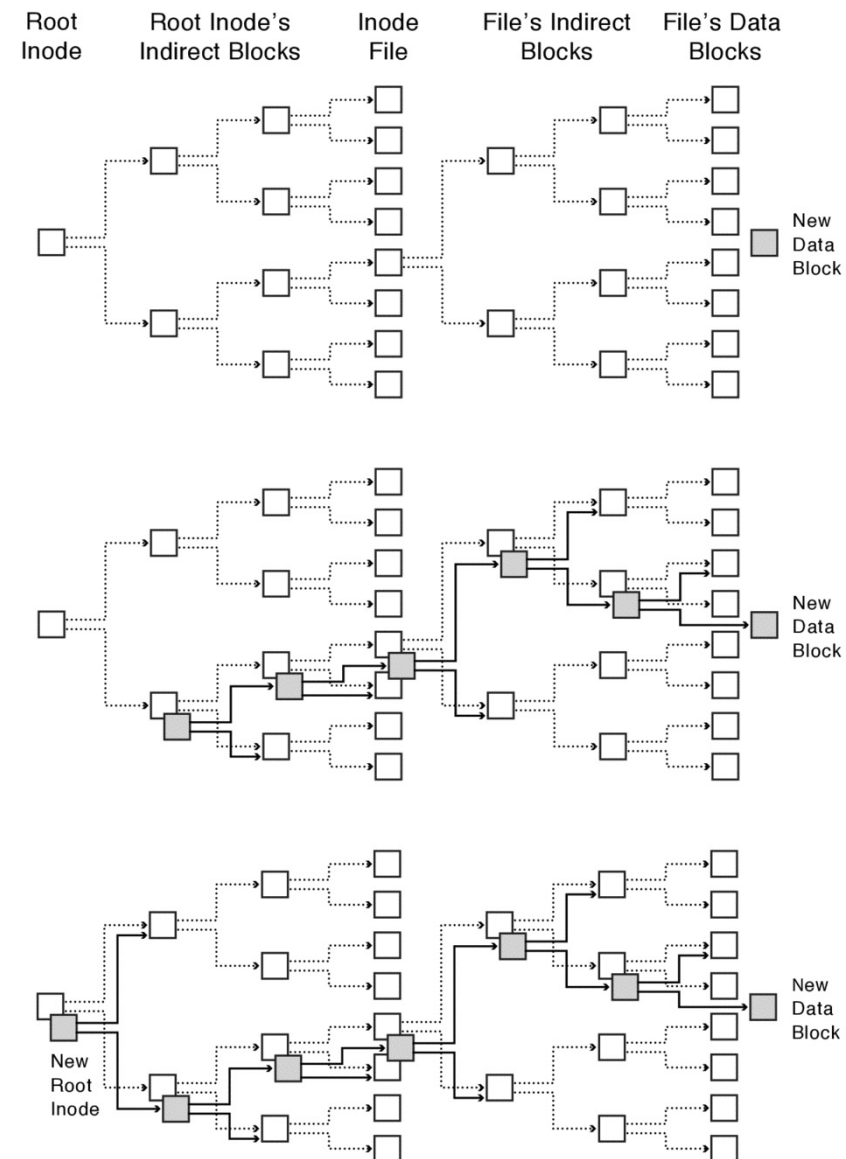
- Two ways to use transactions in file systems: journaling (日志) and logging
- **Journaling**: apply updates to the system's metadata via transactions
 - Microsoft's NTFS, Apple's HFS+, and Linux's XFS/JFS
- **(Full) Logging**: apply both metadata and data in transactions
 - Linux's ext3 and ext4 can be configured to use either journaling or logging

Journaling File Systems

- Applies updates to system metadata (inodes, bitmaps, directories, and indirect blocks) using transactions
 - So those critical data structures are always consistent
- Updates to non-directory files (i.e., user stuff) can be done in place (without logs), full logging optional
 - Avoids writing file contents twice
 - If a program using a journaling file system requires atomic multi-block updates, it needs to provide them itself

Copy-on-Write File System

- To update file system, write a new version of the file system containing the update
 - Never update in place
 - Reuse existing unchanged disk blocks
- Optimization: batch updates
 - Transform many small, random writes into large, sequential writes
- Approach taken in network file server appliances
 - NetApp's Write Anywhere File Layout (WAFL)
 - ZFS (Sun/Oracle) and OpenZFS



Goals for Today


- Transactions for atomic updates
 - Redo Logging
- Redundancy for media failures
 - RAID

RAID: Redundant Arrays of Inexpensive Disks

- Invented by David Patterson, Garth A. Gibson, and Randy Katz here at UCB in 1987
- Data stored on multiple disks (redundancy)
- Either in software or hardware
 - In hardware case, done by disk controller; file system may not even know that there is more than one disk in use
- Initially, five levels of RAID (more now)

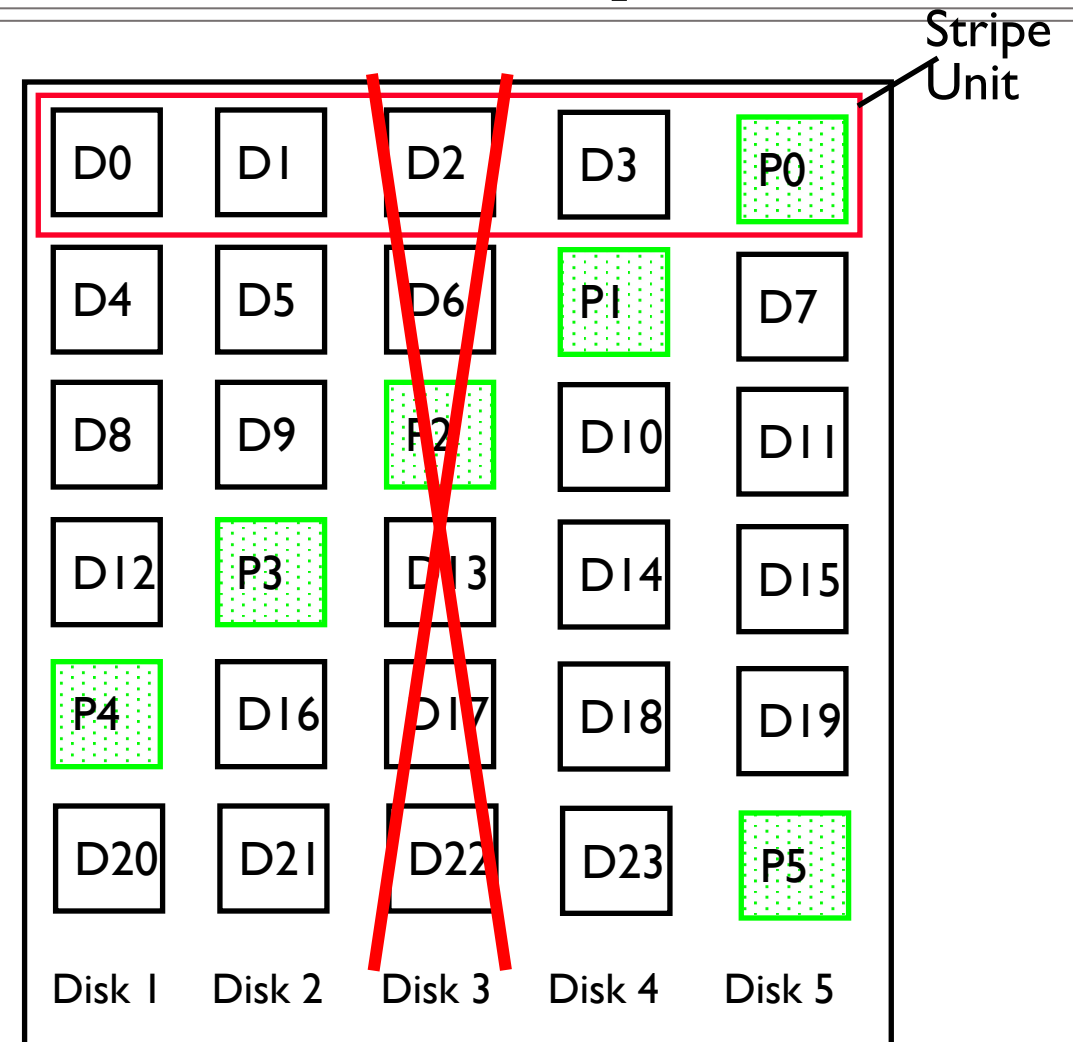
RAID I: Disk Mirroring/Shadowing

- Each disk is fully duplicated onto its “shadow”
 - For high I/O rate, high availability environments
 - Most expensive solution: 100% capacity overhead
- Bandwidth sacrificed on write:
 - Logical write = two physical writes
 - Highest bandwidth when disk heads and rotation fully synchronized (hard to do)
- Reads may be optimized
 - Can have two independent reads to same data
- Recovery:
 - Disk failure \Rightarrow replace disk and copy data to new disk
 - **Hot Spare**: idle disk already attached to system to be used for immediate replacement



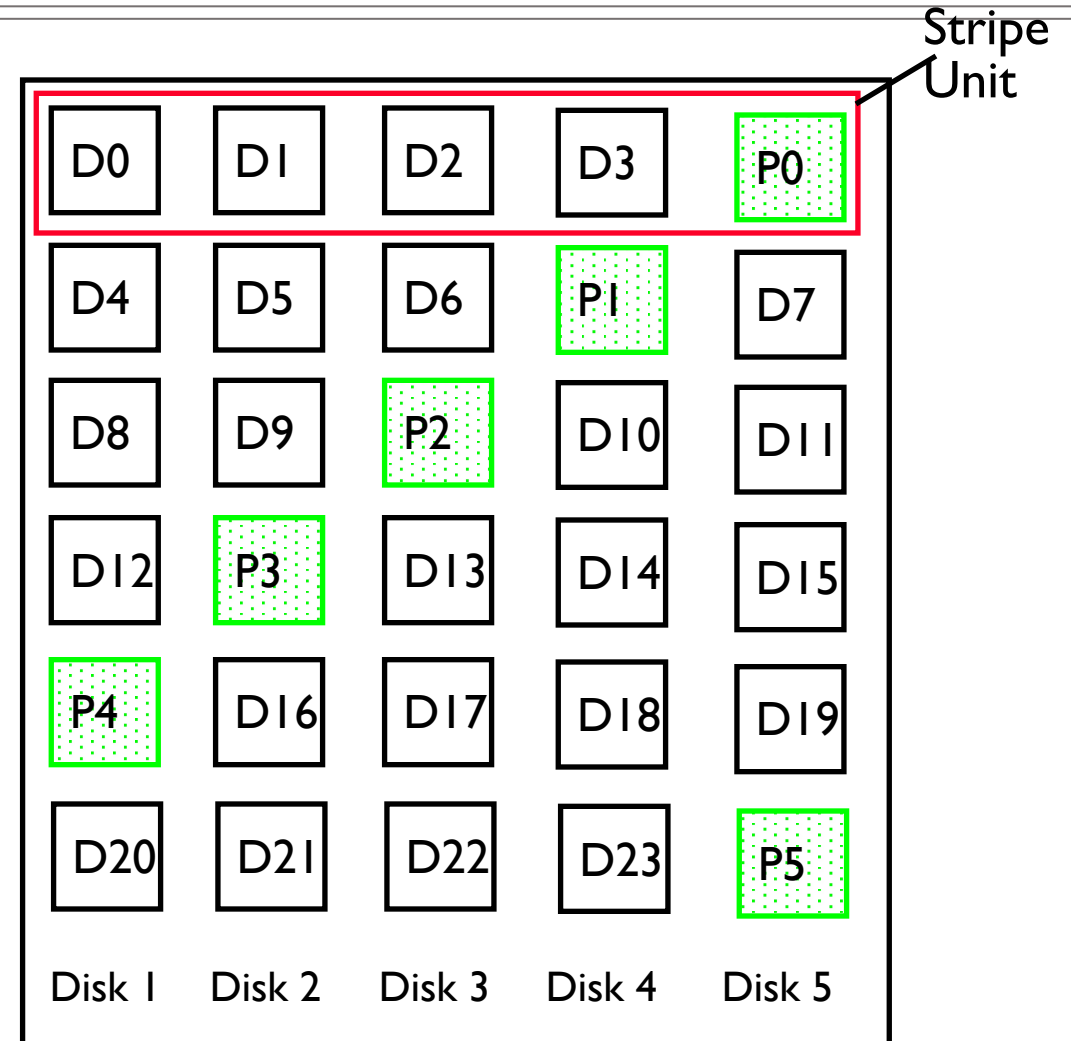
RAID 5+: High I/O Rate Parity

- Data striped across multiple disks
 - Successive blocks stored on successive (non-parity) disks
 - Increased bandwidth over single disk
- Parity block (in green) constructed by XORing (异或) data blocks in stripe
 - $P0 = D0 \oplus D1 \oplus D2 \oplus D3$
 - Can destroy any one disk and still reconstruct data
 - Suppose Disk 3 fails, then can reconstruct: $D2 = D0 \oplus D1 \oplus D3 \oplus P0$




RAID 5+: High I/O Rate Parity

- Rotating parity (奇偶校验)
 - The parity needs to be updated more often than normal data blocks.
- Striping data
 - Balance parallelism vs. sequential access efficiency
- RAID 5 can recover the failed disk only if (i) only one disk fails and (ii) the failed disk is known.




RAID 5+: High I/O Rate Parity

- What I/O operations would occur if we want to update D21 in this figure?



Higher Durability/Reliability through Geographic Replication

- Highly durable – hard to destroy all copies
- Highly available for reads – read any copy
- Low availability for writes
 - Can't write if any one replica is not up
 - Or – need relaxed consistency model
- Reliability? – availability, security, durability, fault-tolerance




Societal Scale Information Systems


- The world is a large distributed system
 - Microprocessors in everything
 - Vast infrastructure behind them



Centralized vs Distributed Systems




Client/Server Model




Peer-to-Peer Model

- **Centralized System:** System in which major functions are performed by a single physical computer
 - Originally, everything on single computer
 - Later: client/server model

Centralized vs Distributed Systems



Client/Server Model



Peer-to-Peer Model

- **Distributed System:** physically separate computers working together on some task
 - Early model: multiple servers working together
 - ❑ Probably in the same room or building
 - ❑ Often called a “cluster”
 - Later models: peer-to-peer/wide-spread collaboration

Distributed Systems: Motivation/Issues/Promise

- Why do we want distributed systems?
 - Cheaper and easier to build lots of simple computers
 - Easier to add power incrementally
 - Users can have complete control over some components
 - Collaboration: much easier for users to collaborate through network resources (such as network file systems)
- The *promise* of distributed systems:
 - Higher availability: one machine goes down, use another
 - Better durability: store data in multiple locations
 - More security: each piece easier to make secure

Distributed Systems: Reality

- Reality has been disappointing
 - Worse availability: depend on every machine being up
 - Lamport: “a distributed system is one where I can’t do work because some machine I’ve never heard of isn’t working!”
 - Worse reliability: can lose data if any machine crashes
 - Worse security: anyone in world can break into system
- Coordination is more difficult
 - Must coordinate multiple copies of shared state information (using only a network)
 - What would be easy in a centralized system becomes a lot more difficult

Distributed Systems: Goals/Requirements

- **Transparency**: the ability of the system to mask its complexity behind a simple interface
- Possible transparencies:
 - **Location**: Can't tell where resources are located
 - **Migration**: Resources may move without the user knowing
 - **Replication**: Can't tell how many copies of resource exist
 - **Concurrency**: Can't tell how many users there are
 - **Parallelism**: System may speed up large jobs by splitting them into smaller pieces
 - **Fault Tolerance**: System may hide various things that go wrong
- Transparency and collaboration require some way for different processors to communicate with one another

Homework- I

- The FastFile file system uses an inode array to organize the files on disk. Each inode consists of a user id (2 bytes), three time stamps (4 bytes each), protection bits (2 bytes), a reference count (2 byte), a file type (2 bytes) and the size (4 bytes). Additionally, the inode contains 13 direct indexes, 1 index to a 1st-level index table, 1 index to a 2nd-level index table, and 1 index to a 3rd level index table. The file system also stores the first 436 bytes of each file in the inode.
 - Assume a disk sector is 512 bytes, and assume that any auxilliary index table takes up an entire sector, what is the maximum size for a file in this system.
 - Is there any benefit for including the first 436 bytes of the file in the inode?

Homework-2

- When user tries to write a file, the file system needs to detect if that file is a directory so that it can restrict writes to maintain the directory's internal consistency. Given a file's name, how would you design a file system to keep track of whether each file is a regular file or a directory?
 - In FAT
 - In FFS
 - In NTFS

Homework-3

- Suppose a variation of FFS includes in each inode 12 direct, 1 indirect, 1 double indirect, 2 triple indirect, and 1 quadruple indirect pointers.
Assuming 6 KB blocks and 6-byte pointers.
 - What is the largest file that can be accessed with direct pointers only?
 - What is the largest file that can be accessed in total?

Homework-4

- Consider a disk queue holding requests to the following cylinders in the listed order: 116, 22, 3, 11, 75, 185, 100, 87. Using the elevator scheduling algorithm, what is the order that the requests are serviced, assuming the disk head is at cylinder 88 and moving upward through the cylinders?

Homework-5

- Search for how different RAID versions (at least 5) work differently and list a table to compare them.